



Apache Spark & Elasticsearch

Holden Karau - Spark Summit 2014

Who am I?

Holden Karau

- Software Engineer @ Databricks
- I've worked with Elasticsearch before
- I prefer she/her for pronouns
- Author of a book on Spark and co-writing another
- github <https://github.com/holdenk>
 - Has all of the code from this talk :)
- e-mail holden@databricks.com
- @holdenkarau

What is Elasticsearch?

- Lucene based distributed search system
- Powerful tokenizing, stemming & other IR tools
- Geographic query support
- Capable of scaling to many nodes

Coffee Berkeley, CA ×

[foursquare Labs, Inc. \[US\] https://foursquare.com/explore?cat=coffee&mode=url&near=Berkeley%2C%20CA](https://foursquare.com/explore?cat=coffee&mode=url&near=Berkeley%2C%20CA)

I'm looking for... Current Map View

Suggestions for Coffee near Berkeley

Search this area

Show me: Specials Haven't Been Friends Price Open Now Saved

1. Asha Tea House
8.7 2086 University Ave (at Shattuck Ave)
Tea Room - \$\$\$\$ - View Menu
2 friends liked this place
"... latte with boba is the bee's knees." (6 tips)
"... tea, but matcha latte (\$3.50) was quite..." (4 tips)
"... else offers hojicha lattes." (2 tips)

Save have been here

2. Philz Coffee
9.2 1600 Shattuck Ave. (at Cedar St.)
Coffee Shop - \$ - \$\$\$
Josh A. liked this place
"... to your Iced Mocha Tesora." (7 tips)
"The mint mojito iced coffee is maybe the best..." (5 tips)
"... Mojito and Gingersnap ice coffee are the nest..." (3 tips)

Save +2

3. Caffe Strada
8.9 2300 College Ave (at Bancroft Way)
Coffee Shop - \$ - \$\$\$
"grab a strada bianca mocha and hang out on the..." (3 tips)
"... Bianca Mocha + nut bun" (2 tips)
"... mac, but it wouldn't work, other people were having..." (2 tips)

Save have been here

4. Starbucks
7.6 2224 Shattuck Ave
Coffee Shop - \$\$\$\$ - View Menu
Starbucks - May 21, 2014
Looking to try something unique? Try one of our Reserve Coffees. Our Brazil Bourbon Rio Verde is a beautifully balanced cup with notes of plum and chocolate.
People talk about: venti, peppermint

Save cooper b. has been here.

5. Au Coquelet Cafe
8.0 2000 University Ave (at Milvia St)
Cafe - \$\$\$\$ - View Menu
"... sure you charge your laptop before you get there -..." (2 tips)
"Open late, excellent study/drawing..." (3 tips)
"... Mousse Torte. Worth going out of your way..." (2 tips)

Save have been here

6. Uncommon Cafe
5.9 2813 7th St
Coffee Shop - \$\$\$\$ - View Menu
Cass M. liked this place

The map displays the city of Berkeley, California, with various neighborhoods labeled: ANNEX, EAST SHORE STATE PARK, ALBANY BULB, ALBANY HILL PARK, SUNSET VIEW CEMETERY, THOUSAND OAKS, CRAGMONT, BERKELEY HILLS, TILDEN PARK, LA LOMA PARK, TILDEN GOLF COURSE, NORTHBRAE, UNIVERSITY VILLAGE, WESTBRAE, NORTHSIDE, BERKELEY, SOUTHSIDE, PANORAMIC HILL, CLAREMONT, ELMWOOD, ROCKRIDGE, TEMSCAL, CLAREMONT COUNTRY CLUB, MOUNTAIN VIEW CEMETERY, and PIEDMONT AVENUE. The map also shows the San Pablo Dam, Grizzly Peak Blvd, and several state routes (CA 13, CA 24, Tunnel Rd, Grove Shafter Rd, Warren Pkwy). Numerous coffee shop locations are marked with blue pins, each numbered from 1 to 29. The pins are concentrated in the central Berkeley area, along Shattuck Ave, University Ave, and surrounding streets like Gilman, Berkeley, and Dwight Way.

Talk overview

Goal: understand how to work with ES & Spark

- Spark & Spark streaming let us re-use indexing code
- We can customize the ES connector to write to the shard based on partition
- Illustrate with twitter & show top tags per region
- Maybe a live demo of the above demo*

Assumptions:

- Familiar(ish) with Search
- Can read Scala

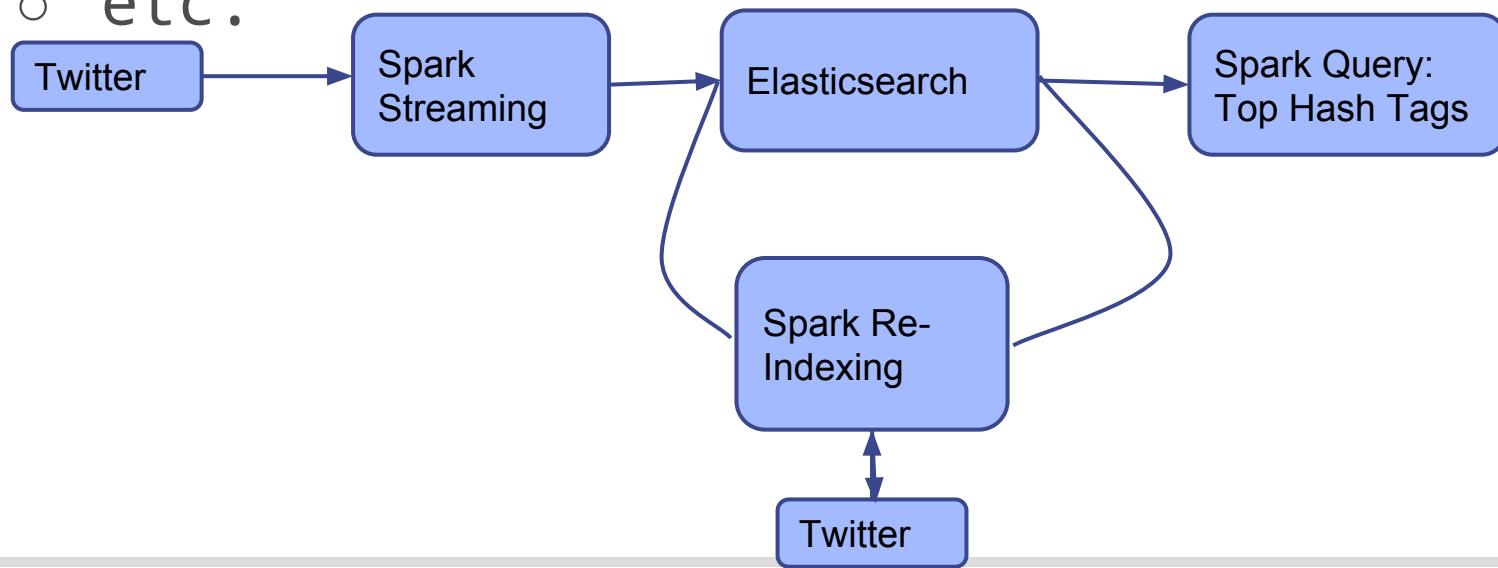
Things you don't have to worry about:

- All the code is on-line, so don't worry if you miss some

*If we have extra time at the end

Spark + Elasticsearch

- We can index our data on-line & off-line
- Gain the power to query our data
 - based on location
 - free text search
 - etc.



Why should you care?

Small differences between off-line and on-line

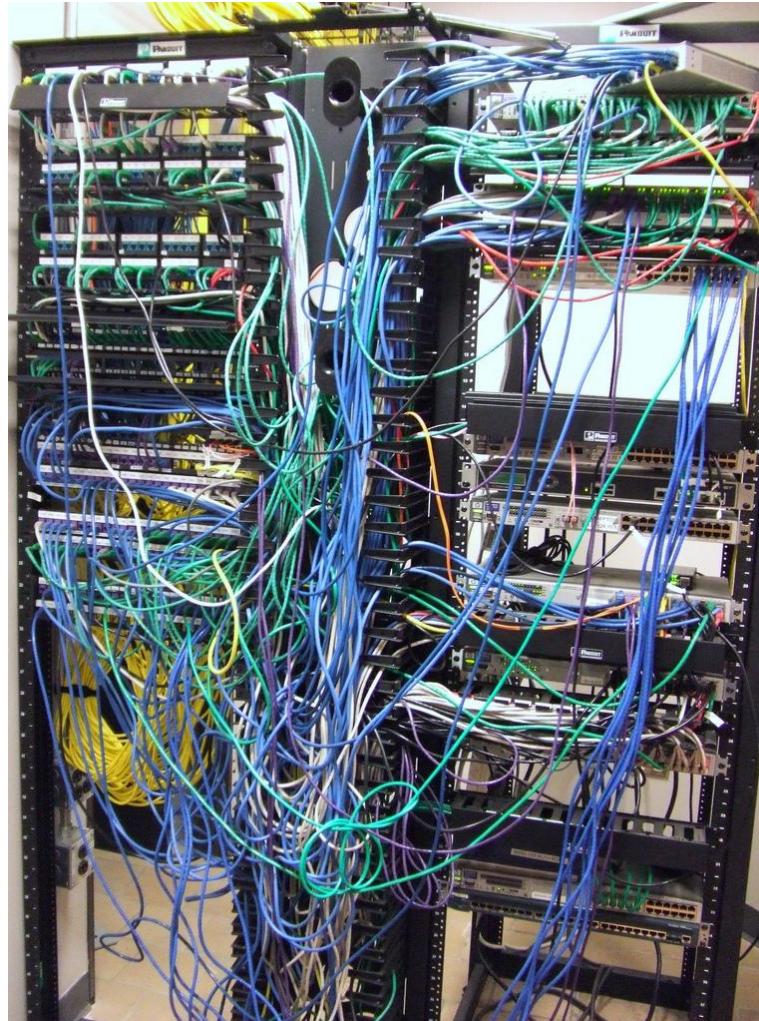




Lets start with the on-line pipeline

```
val ssc = new StreamingContext(master, "IndexTweetsLive",  
Seconds(1))  
  
val tweets = TwitterUtils.createStream(ssc, None)
```

Lets get ready to write the data into Elasticsearch



Lets get ready to write the data into Elasticsearch

```
def setupEsOnSparkContext(sc: SparkContext) = {  
    val jobConf = new JobConf(sc.hadoopConfiguration)  
  
    jobConf.set("mapred.output.format.class",  
               "org.elasticsearch.hadoop.mr.EsOutputFormat")  
  
    jobConf.setOutputCommitter(classOf[FileOutputCommitter])  
  
    jobConf.set(ConfigurationOptions.ES_RESOURCE_WRITE,  
               "twitter/tweet")  
  
    FileOutputFormat.setOutputPath(jobConf, new Path("-"))  
    jobconf  
}
```

Add a schema

```
curl -XPUT 'http://localhost:  
9200/twitter/tweet/_mapping' -d '  
{  
  "tweet" : {  
    "properties" : {  
      "message" : {"type" : "string"},  
      "hashTags" : {"type" : "string"},  
      "location" : {"type" : "geo_point"}  
    }  
  }  
}  
'
```

Lets format our tweets

```
def prepareTweets(tweet: twitter4j.Status) = {  
  ...  
  val hashTags = tweet.getHashtagEntities().map(_.getText())  
  HashMap(  
    "docid" -> tweet.getId().toString,  
    "message" -> tweet.getText(),  
    "hashTags" -> hashTags.mkString(" "),  
    "location" -> s"$lat,$lon"  
  )  
}  
}  
// Convert to HadoopWritable types  
mapToOutput(fields)  
}
```

And save them...

```
tweets.foreachRDD{(tweetRDD, time) =>
  val sc = tweetRDD.context
  // The jobConf isn't serializable so we create it here
  val jobConf = SharedESConfig.setupEsOnSparkContext(sc,
    esResource, Some(esNodes))
  // Convert our tweets to something that can be indexed
  val tweetsAsMap = tweetRDD.map(
    SharedIndex.prepareTweets)
  tweetsAsMap.saveAsHadoopDataset(jobConf)
}
```

Now let's query them!

```
{"filtered" : {  
    "query" : {  
        "match_all" : {}  
    }  
, "filter" :  
    {"geo_distance" :  
        {  
            "distance" : "${dist}km",  
            "location" :  
                {  
                    "lat" : "${lat}",  
                    "lon" : "${lon}"  
                }  
        }  
    }  
}
```

Now let's find the hash tags :)

```
// Set our query
jobConf.set("es.query", query)

// Create an RDD of the tweets
val currentTweets = sc.hadoopRDD(jobConf,
    classOf[EsInputFormat[Object, MapWritable]],
    classOf[Object], classOf[MapWritable])

// Convert to a format we can work with
val tweets = currentTweets.map{ case (key, value) =>
    SharedIndex.mapWritableToInput(value) }

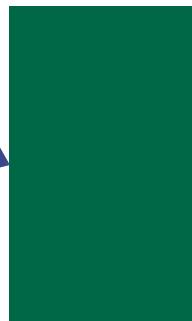
// Extract the hashtags
val hashTags = tweets.flatMap{t =>
    t.getOrElse("hashTags", "").split(" ")}
```

and extract the top hashtags

```
object WordCountOrdering extends Ordering[(String, Int)]{  
    def compare(a: (String, Int), b: (String, Int)) = {  
        b._2 compare a._2  
    }  
}  
  
val ht = hashtags.map(x => (x, 1)).reduceByKey((x,y) => x+y)  
  
val topTags = ht.takeOrdered(40)(WordCountOrdering)
```

NYC

- #MEX,11
- #Job,11
- #Jobs,8
- #nyc,7
- #CarterFollowMe,7
- #Mexico,6
- #BRA,6
- #selfie,5
- #TweetMyJobs,5
- #LHHATL,5
- #NYC,5
- #ETnow,4
- #TeenWolf,4
- #CRO,4



SF

- #Job,6
- #Jobs,5
- #MEX,4
- #TweetMyJobs,3
- #TeenWolfSeason4,2
- #CRO,2
- #Roseville,2
- #Healthcare,2
- #GOT7COMEBACK,2
- #autodeskinterns,2



Indexing Part 2

(electric boogaloo)

The screenshot shows the Elasticsearch Head plugin interface at <http://192.168.7.8:9200/>. The top bar displays "ElasticSearch" and the URL. The status is "Rick cluster health: yellow (6, 18)". Below the header, there are tabs for "Cluster Overview" and "New Index".

The main area shows several indices with their sizes and document counts:

- cu_docs**: size: 180Gb (540Gb), docs: 995131 (995131). Shards: 0 (green) and 1 (green).
- bnil**: size: 80kb (480kb), docs: 90 (90). Shards: 0 (green) and 1 (green).
- cu_msg**: size: 313Gb (1.56Tb), docs: 10047450 (10140915). Shards: 0 (green) and 1 (green). A context menu is open over shard 1, showing options: Refresh, Flush, Gateway Snapshot, Test Analyser, Close, and Delete... The "Gateway Snapshot" option is highlighted.
- anvil**: index: close. Shards: 0 (green) and 1 (green).

Below these, individual nodes are listed with their names, IDs, and IP addresses:

- Leon**: 3Wqr1xaCRu-b0uEzDkmrDg, inet[/192.168.7.8:9202]. Shards: 0 (green) and 1 (green).
- Pris**: L8qx7lfSI-kcKq_6bMbWw, inet[/192.168.7.8:9204]. Shards: 0 (green) and 1 (green).
- Rick**: Vnpri1FNTGirwRfZsZ2RxQ, inet[/192.168.7.8:9200]. Shards: 1 (green) and 2 (green).
- Rachel**: 87KsIv0FTVskkqwENaja6A, inet[/192.168.7.8:9203]. Shards: 1 (green) and 2 (green).
- Zhora**: b6NxRTxsR_WUQj5cXPKhbw, inet[/192.168.7.8:9205]. Shards: 0 (green) and 2 (green).
- Roy**: _8Rl2wYVT7Svn_v5F97jJA, inet[/192.168.7.8:9201]. Shards: 0 (green) and 2 (green).
- Unassigned**: Shards: 0 (grey) and 1 (grey).

A modal window is open over node "Rick" (shard 2), displaying its details:

```
{  
  name: "Leon",  
  transport_address: "inet[/192.168.7.8:9302]",  
  attributes: {},  
  http_address: "inet[/192.168.7.8:9202]",  
  os:  
    refresh_interval: 5000,  
    cpu:  
      vendor: "Intel",  
      model: "Macmini4,1",  
      mhz: 2400,  
      total_cores: 2,  
      total_sockets: 1,  
      cores_per_socket: 2,  
      cache_size: "3kb",  
      cache_size_in_bytes: 3072  
}
```

Writing directly to a node with the correct shard saves us network overhead

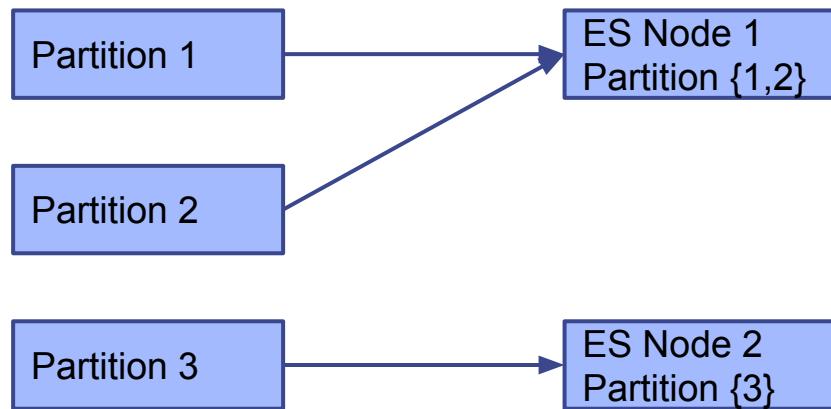
So what does that give us?

Spark sets the filename to part-[partition number]

We can modify the connector to parse this! :D

If we have same partitioner we write directly

Likely the best place to use this is in re-indexing data



Re-index all the things*

```
// Fetch them from twitter
val t4jt = tweets.flatMap{ tweet =>
    val twitter = TwitterFactory.getSingleton()
    val tweetID = tweet.getOrElse("docid", "")
    Option(twitter.showStatus(tweetID.toLong))
}

t4jt.map(SharedIndex.prepareTweets)
    .saveAsHadoopDataset(jobConf)
```

“Useful” links

- Feedback: holden@databricks.com
- Customized ES connector*: <https://github.com/holdenk/elasticsearch-hadoop>
- Demo code: <https://github.com/holdenk/elasticsearchspark>
- Elasticsearch: <http://www.elasticsearch.org/>
- Spark: <http://spark.apache.org/>
- Spark streaming: <http://spark.apache.org/streaming/>
- Elasticsearch Spark documentation: <http://www.elasticsearch.org/guide/en/elasticsearch/hadoop/current/spark.html>

So what did we cover?

- Indexing data with Spark to Elasticsearch
- Sharing indexing code between Spark & Spark Streaming
- Using Elasticsearch for geolocal data in Spark
- Making our indexing aware of Elasticsearch
- Lots* of cat pictures



Cat photo from <https://www.flickr.com/photos/deerwooduk/579761138/in/photolist-4GCc4z-4GCbAV-6Ls27-34evHS-5UBnJv-TeqMG-4iNNn5-4w7s61-6GMLYS-6H5QWY-6aJLUT-tqfrf-6mJ1Lr-84kGX-6mJ1GB-vVqN6-dY8aj5-y3jK-7C7P8Z-azEtd/>